

ASSIGNMENT-04

Name: Khushi Sharma

Roll no: 2301410007

Course: B-Tech CSE (Cyber Security)

i) imagine two people trying to withdraw money from the same bank account at the same ATM at the same time. Both see the balance as 500. Each withdraws ₹500, & the system allows both, causing inconsistency. The final outcome depends on *who finishes first* not on correct logic → This is a race condition

→ How MUTUAL EXCLUSION SOLVES IT :-

A mutex would allow only one user to access / update the account at a time. When one withdrawal is in progress, the other must wait. This prevents conflicting updates & ensures consistent balance.

2) Peterson's Solution:-

- Pure software - level algorithm; simple but error-prone; limited to 2 processes.
- Requires strict assumptions about atomic read/write ordering → not reliable on modern multi-core CPUs.

→ Semaphore :-

- More General :- easier for programmers ; widely supported.
- Depends on hardware atomic operations like test-and-set ; explicitly supported by processors.

③ Monitors automatically combine mutual exclusion with condition variables , so synchronization is handled implicitly by the language / runtime.

Advantage :- They prevent low-level race-conditions automatically , allowing threads on multiple cores to safely access shared data without manually managing semaphores , reducing bugs & complexity.

④ Starvation :- If the system gives priority to readers a continuous stream of readers can keep arriving , causing a writer waiting until to have get access.

→ Prevention Method :- use writer preference
 and a writer requests the lock, no new readers are allowed to enter. Existing readers finish, then the writer is served.
 This ensures no indefinite waiting.

- (5) To prevent deadlock, processes must request all resources at once before starting.

→ Practical Drawback :-

This leads to low resource utilization.
 A process may hold several resources it is not currently using, while others wait unnecessarily. This increases waiting time & reduces system throughput.

- (6) a) Combining all edges from three edge fragments graphically :-

$$P_1 \rightarrow P_2 \rightarrow P_3 \rightarrow P_6$$

$$\text{& separately :- } P_3 \rightarrow P_4$$

- b) Detect deadlock & list processes involved
 → A deadlock exists if there is a cycle in the wait-for-graph
 → we have the cycle $P_1 \rightarrow P_2 \rightarrow P_5 \rightarrow P_6 \rightarrow P_1$, processes involved in the deadlock $\rightarrow P_1, P_2, P_5, P_6$.

(c) Suggest one distributed deadlock detection.

→ chand - Miru - Mass (probe) algorithm is started choice for distributed deadlock detection.

How it works :-

- A process that suspect it is blocked initiates probe messages of form :-

(4) Given :-

- Local access time $T_L \geq 5\text{ ms}$
- Remote access time $T_R \geq 25\text{ ms}$
- Probability of file is remote
 $= p \rightarrow 0.3$

a) Expected time $E[T]$ is :-

$$E[T] = p \cdot T_R + (1-p) \cdot T_L$$

Plug Numbers :-

$$E[T] = 0.3 \times 25 + 0.7 \times 5 = 7.5 + 3.5$$

$$= 11.0\text{ ms}$$

b) Caching strategy suggestion :-

→ suggested strategy :- client site LRU caching with write-back for read-heavy file & validation on TTL.

Justification :-

- LRU (Least Recently Used) :- works well in practice because file access pattern often has temporal locality or recently used file are likely to be reused.

(f)

a) proposed optimal minimum steps :-

- Take one full checkpoint every 10s.
- Take incremental checkpoint every 3 b/w full checkpoints.

Total checkpoint overhead periods :-

- Full : $1 \times 200 \text{ ms} = 20 \text{ ms}$
- incrementals : $9 \times 50 \text{ ms} = 450 \text{ ms}$
Total $\geq 650 \text{ ms per 10s}$
- average overhead :- 65 ms/c

b) Reasoning :-

- RPO constraint : The system must be prepared to restore a state no older than α ; therefore an incremental checkpoint must be taken at least once per second.
- Full checkpoint all expensive (200 ms). Doing them less frequently reduces heavy overhead full checkpoint reduces recovery time bcoz only one full.
- Trade-off :- Frequent incremental add modest overhead but keep RPO tight! Periodic full keeps recovery efficient.

⑨

a) KEY CHALLENGES :-

- i) Sudden bursty traffic :- Requests spike orders of magnitude within seconds.
- ii) Geographic latency & data locality :- User should be served from the nearest region when possible.
- iii) Hot-spots / skewed load :- certain

products / regions get disproportion-
ate attention.

iV) Fairness & SLA guarantees :- Mission -
critical request (Payment) must get
prioritized.

b.) Fault - tolerance strategy :-

1. use multi-region Deployment.
2. Replicate data based on RPO Requirements.
3. Enable Automated failover to meet RTO.
4. Maintain Cross-Region backup & regular PR drills.